



Applications

Our Goals:

- conceptual, implementation aspects of network application protocols
 - transport-layer service models
 - client-server paradigm
 - peer-to-peer paradigm
 - content distribution networks
- learn about protocols by examining popular application-level protocols
 - HTTP
 - FTP
 - SMTP / POP3 / IMAP
 - DNS
- creating network applications
 - socket API



<https://www.youtube.com/watch?v=uGI00HV7Cfw>



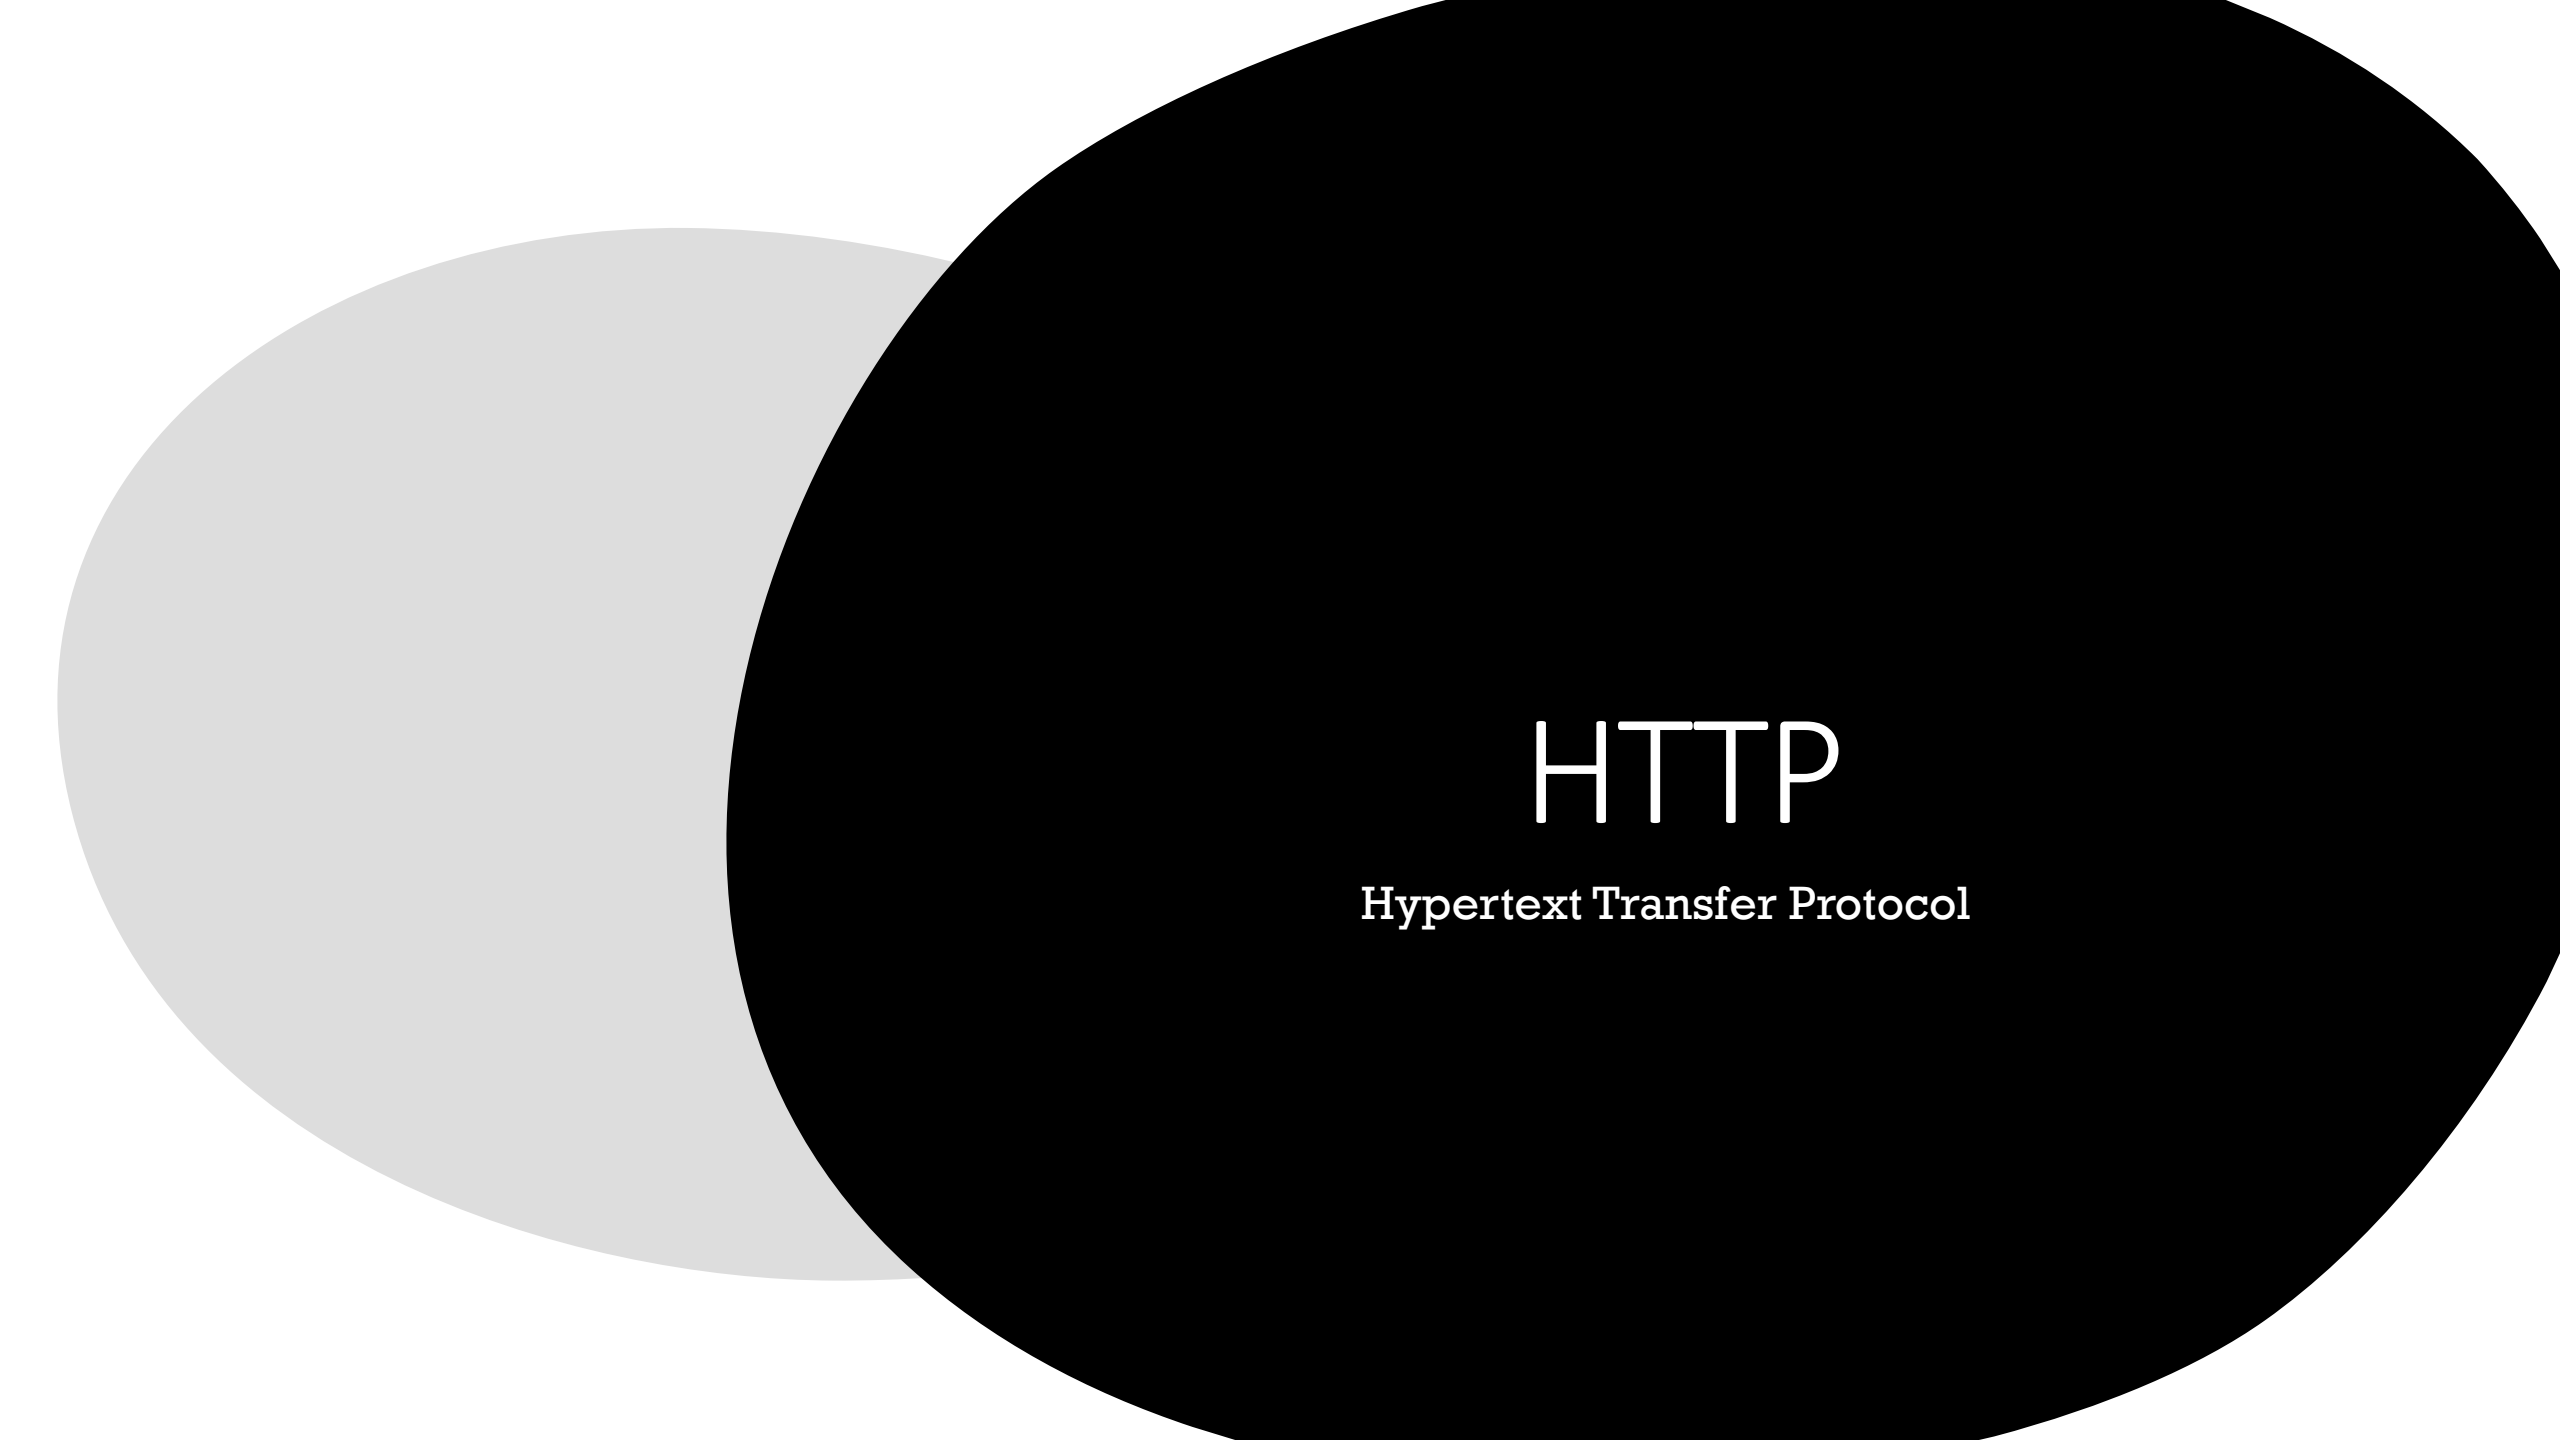
The evolution of a computing technology

Computing Era

Institutional computing
One device for many users

Mass market
One device for one user

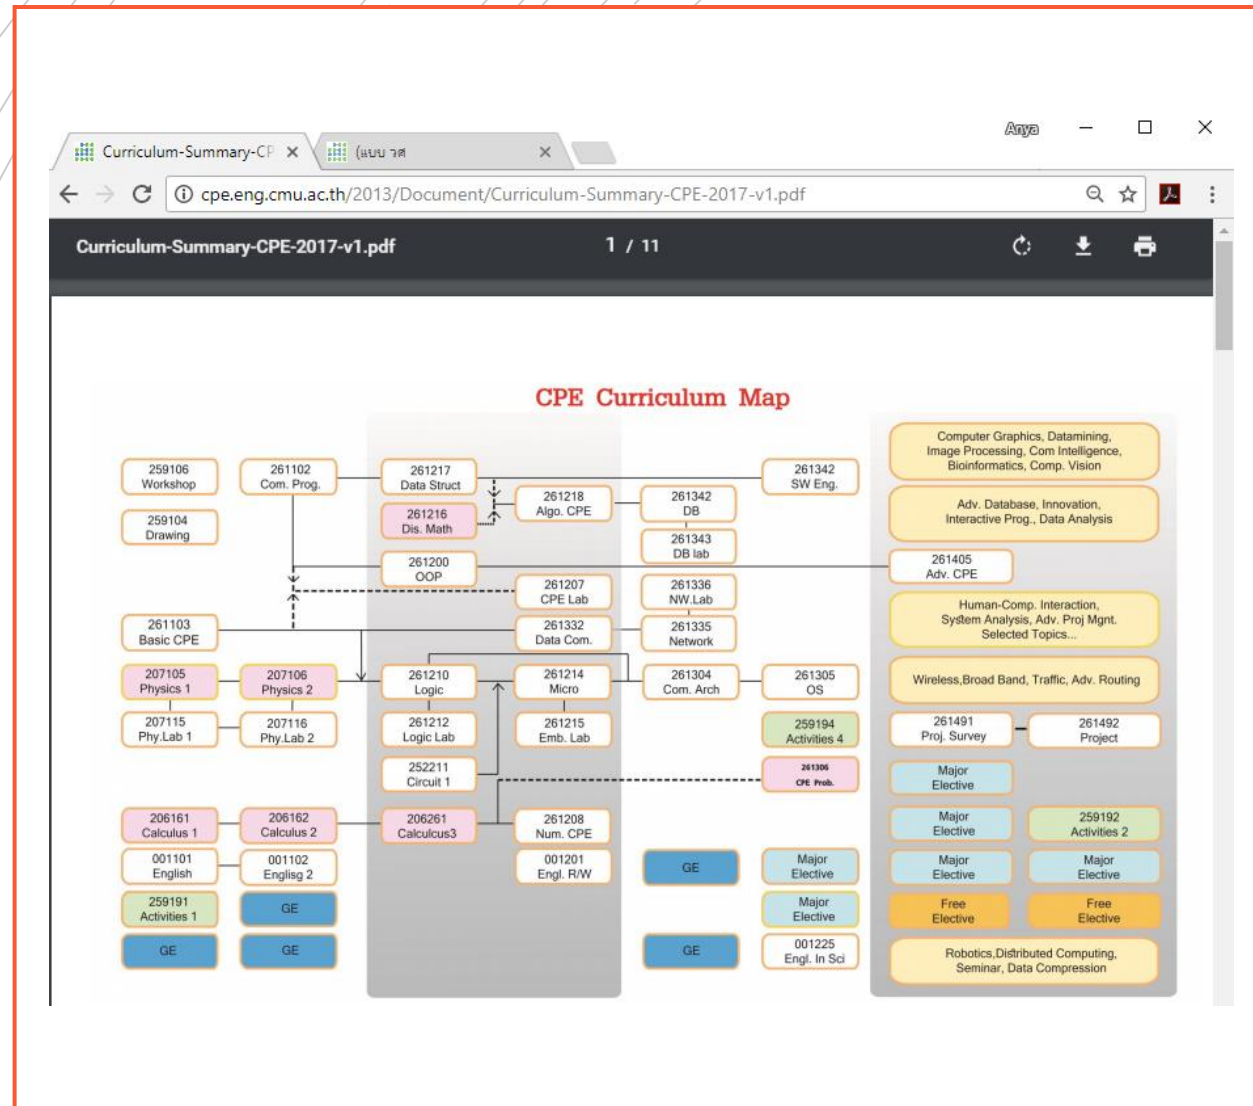
Multi-device computing
Many devices for one user

The image features two large, overlapping circles. The circle on the left is a light gray, and the circle on the right is black. The black circle is positioned slightly to the right and overlaps the gray circle. Centered within the black circle is the text 'HTTP' in a white, serif font.

HTTP

Hypertext Transfer Protocol

HOW TO FETCH A WEBPAGE?



- When user type a URL...
- URL indicate hostname and path to “objects”
- Interaction between Server/Client

HTTP overview



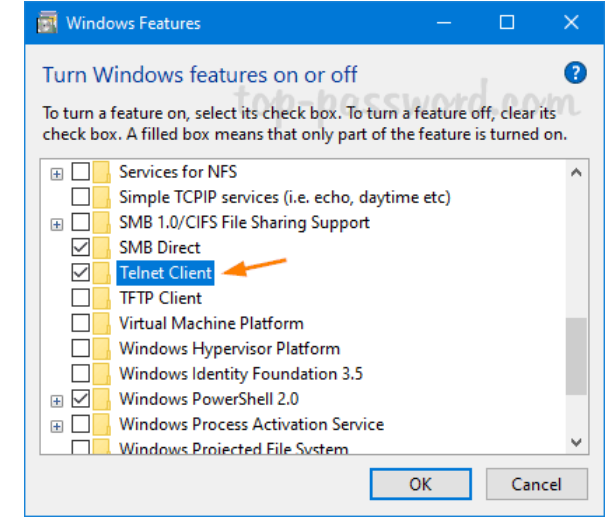
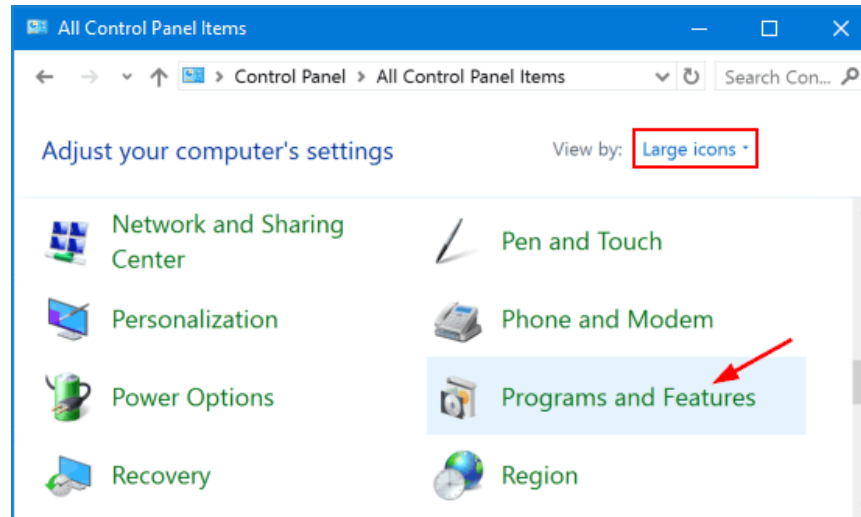
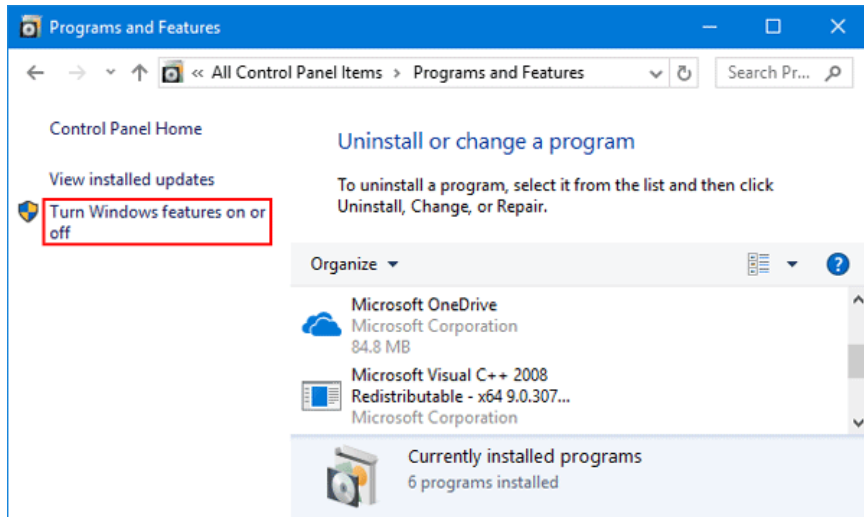
HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - **server**: Web server sends (using HTTP protocol) objects in response to requests

Web and HTTP

First, a review...

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,



Telnet on Windows : Enable Telnet

- Open Terminal with cmd command
- Enable Telnet on Windows
 - Open "Control Panel > Programs" and select "Turn Windows features on or off" under the category "Programs and Features".
 - In the Windows Features list, select "Telnet Client" to activate it, then apply the change.

- Terminal
- Keyboard
- Bell
- Features
- Window
- Appearance
- Behaviour
- Translation
- + Selection
- Colours
- Connection
- Data
- Proxy
- SSH

Basic options for ,

Specify the destination you want to connect to:

Host Name (or IP address)
cpe.eng.cmu.ac.th 80

Connection type:
☐ SSH ☐ Serial ☒ Other: Raw

Load, save or delete a stored session

Saved Sessions

Default Settings	Load
TelnetHTTPSMTP	

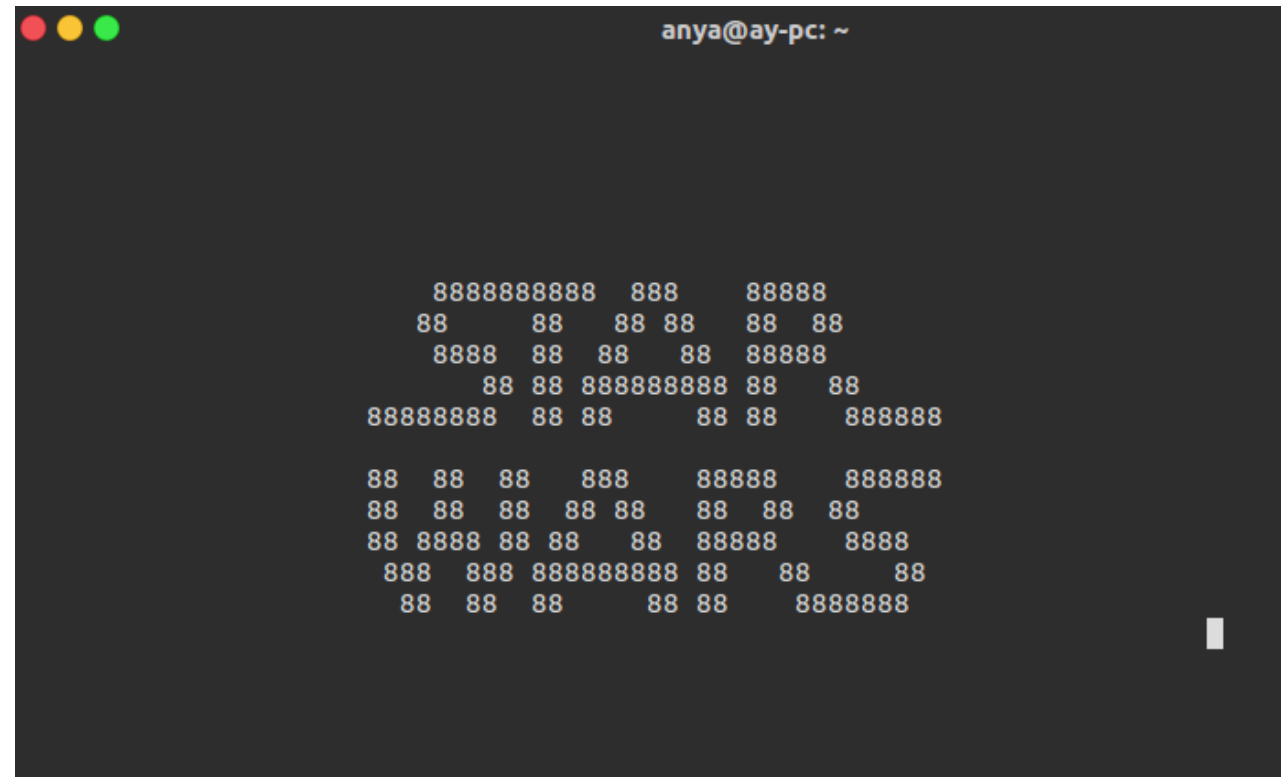
Save

Delete

Telnet with
Putty

```
telnet towel.blinkenlights.nl 23 ?
```

Telnet



Telnet provides a bidirectional interactive text-oriented over the Transmission Control Protocol (TCP).

Results from TELNET

```
anya@ay-pc:~$ telnet www.cpe.eng.cmu.ac.th 80
```

```
Trying 202.28.24.139...
```

```
Connected to cpe.eng.cmu.ac.th.
```

```
Escape character is '^['.
```

```
GET / HTTP/1.1
```

```
Host: www.cpe.eng.cmu.ac.th
```

```
GET / HTTP/1.1
```

```
Host: www.cpe.eng.cmu.ac.th
```

```
HTTP/1.1 200 OK
```

```
Date: Tue, 10 Oct 2017 07:00:26 GMT
```

```
Server: Apache
```

```
Last-Modified: Tue, 22 Oct 2013 06:07:54 GMT
```

```
ETag: "5ea04a-187-4e94e38a68280"
```

```
Accept-Ranges: bytes
```

```
Content-Length: 391
```

```
Content-Type: text/html
```

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN" "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
```

```
<html xmlns="http://www.w3.org/1999/xhtml">
```

```
<head>
```

```
<meta http-equiv="Content-Type" content="text/html; charset=utf-8" />
```

```
<META HTTP-EQUIV="Refresh" CONTENT="0;URL=http://cpe.eng.cmu.ac.th/2013">
```

```
<title>Computer Engineering CMU</title>
```

```
</head>
```

```
<body>
```

```
</body>
```

```
</html>
```

```
Connection closed by foreign host.
```

```
anya@ay-pc:~$
```



http ✕ ➡ Expression... +

No.	Time	Source	Destination	Protocol	Length	Info
1185	3.546610112	10.10.182.50	128.119.245.12	HTTP	431	GET /wireshark-labs/HTTP-wireshark-file1.
1435	4.142205464	128.119.245.12	10.10.182.50	HTTP	552	HTTP/1.1 200 OK (text/html)
1437	4.161056568	10.10.182.50	128.119.245.12	HTTP	372	GET /favicon.ico HTTP/1.1
1616	4.460635821	128.119.245.12	10.10.182.50	HTTP	550	HTTP/1.1 404 Not Found (text/html)

- ▶ Frame 1185: 431 bytes on wire (3448 bits), 431 bytes captured (3448 bits) on interface 0
- ▶ Ethernet II, Src: Micro-St_26:88:10 (d8:cb:8a:26:88:10), Dst: Industri_29:95:1b (00:00:cd:29:95:1b)
- ▶ Internet Protocol Version 4, Src: 10.10.182.50, Dst: 128.119.245.12
- ▶ Transmission Control Protocol, Src Port: 55544, Dst Port: 80, Seq: 1, Ack: 1, Len: 365
- ▼ Hypertext Transfer Protocol
 - ▶ GET /wireshark-labs/HTTP-wireshark-file1.html HTTP/1.1\r\n
 - Host: gaia.cs.umass.edu\r\n
 - User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:55.0) Gecko/20100101 Firefox/55.0\r\n
 - Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8\r\n
 - Accept-Language: en-US,en;q=0.5\r\n
 - Accept-Encoding: gzip, deflate\r\n
 - Connection: keep-alive\r\n
 - Upgrade-Insecure-Requests: 1\r\n
 - \r\n
 - [Full request URI: http://gaia.cs.umass.edu/wireshark-labs/HTTP-wireshark-file1.html]
 - [HTTP request 1/2]
 - [Response in frame: 1435]
 - [Next request in frame: 1437]

Results from WIRESHARK
REQUEST

Wireshark interface showing packet capture on *enp3s0. The filter is set to http.

No.	Time	Source	Destination	Protocol	Length	Info
1185	3.546610112	10.10.182.50	128.119.245.12	HTTP	431	GET /wireshark-labs/HTTP-wireshark-fi
1435	4.142205464	128.119.245.12	10.10.182.50	HTTP	552	HTTP/1.1 200 OK (text/html)
1437	4.161056568	10.10.182.50	128.119.245.12	HTTP	372	GET /favicon.ico HTTP/1.1
1616	4.460635821	128.119.245.12	10.10.182.50	HTTP	550	HTTP/1.1 404 Not Found (text/html)

Frame 1435: 552 bytes on wire (4416 bits), 552 bytes captured (4416 bits) on interface 0

Ethernet II, Src: Industri_29:95:1b (00:00:cd:29:95:1b), Dst: Micro-St_26:88:10 (d8:cb:8a:26:88:10)

Internet Protocol Version 4, Src: 128.119.245.12, Dst: 10.10.182.50

Transmission Control Protocol, Src Port: 80, Dst Port: 55544, Seq: 1, Ack: 366, Len: 486

Hypertext Transfer Protocol

- HTTP/1.1 200 OK\r\n
 Date: Mon, 16 Oct 2017 10:33:49 GMT\r\n
 Server: Apache/2.4.6 (CentOS) OpenSSL/1.0.2k-fips PHP/5.4.16 mod_perl/2.0.10 Perl/v5.16.3\r\n
 Last-Modified: Mon, 16 Oct 2017 05:59:01 GMT\r\n
 ETag: "80-55ba3b4031322"\r\n
 Accept-Ranges: bytes\r\n
 Content-Length: 128\r\n
 Keep-Alive: timeout=5, max=100\r\n
 Connection: Keep-Alive\r\n
 Content-Type: text/html; charset=UTF-8\r\n
 \r\n
 [HTTP response 1/2]
 [Time since request: 0.595595352 seconds]
 [Request in frame: 1185]
 [Next request in frame: 1437]
 [Next response in frame: 1616]
 File Data: 128 bytes

Results from WIRESHARK
RESPONSE

HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**

- ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

header
lines

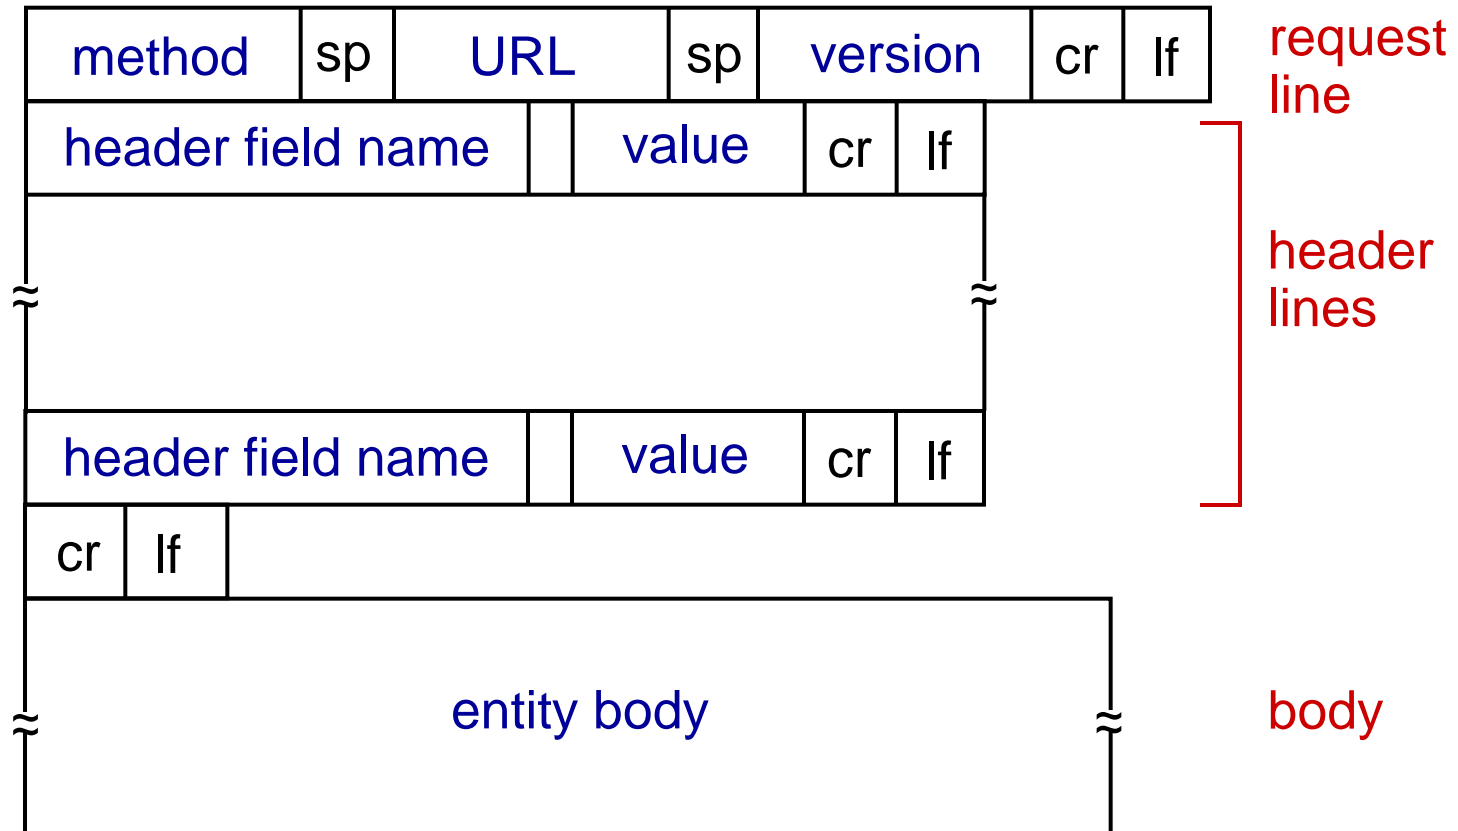
carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

* Check out the online interactive exercises for more
examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

HTTP request message: general format



Method types

HTTP/1.0:

- GET
- POST
- HEAD
 - asks server to leave requested object out of response

HTTP/1.1:

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field

HTTP response message

status line
(protocol
status code
status phrase)

header
lines

data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
1\r\n
\r\n
data data data data data ...
```

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this msg

301 Moved Permanently

- requested object moved, new location specified later in this msg
(Location:)

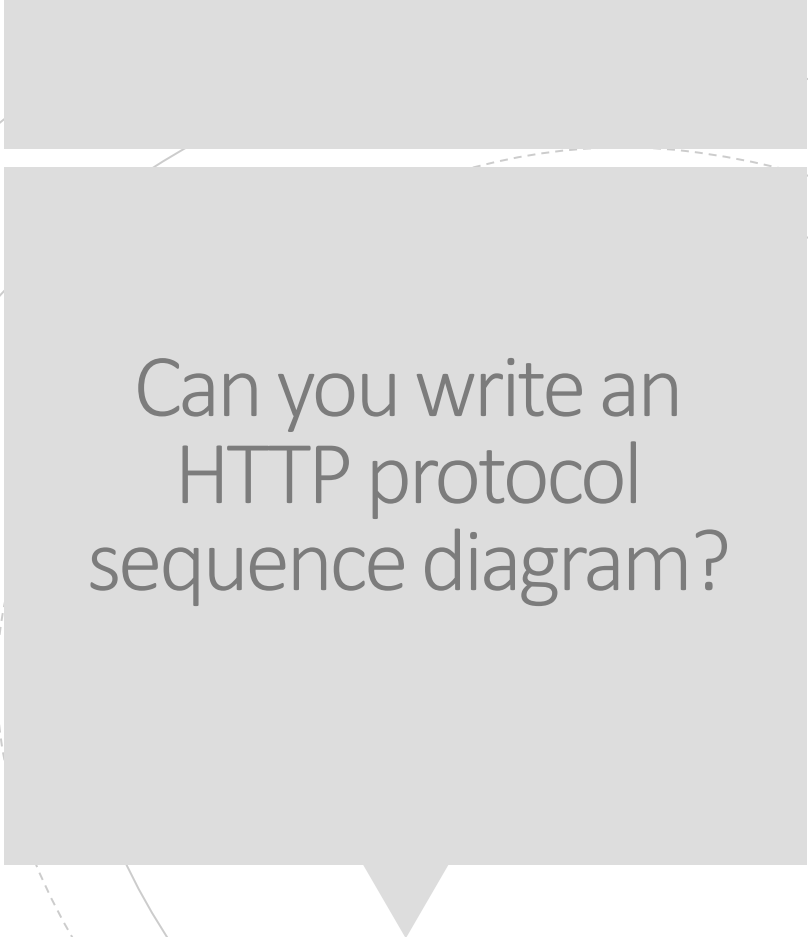
400 Bad Request

- request msg not understood by server

404 Not Found

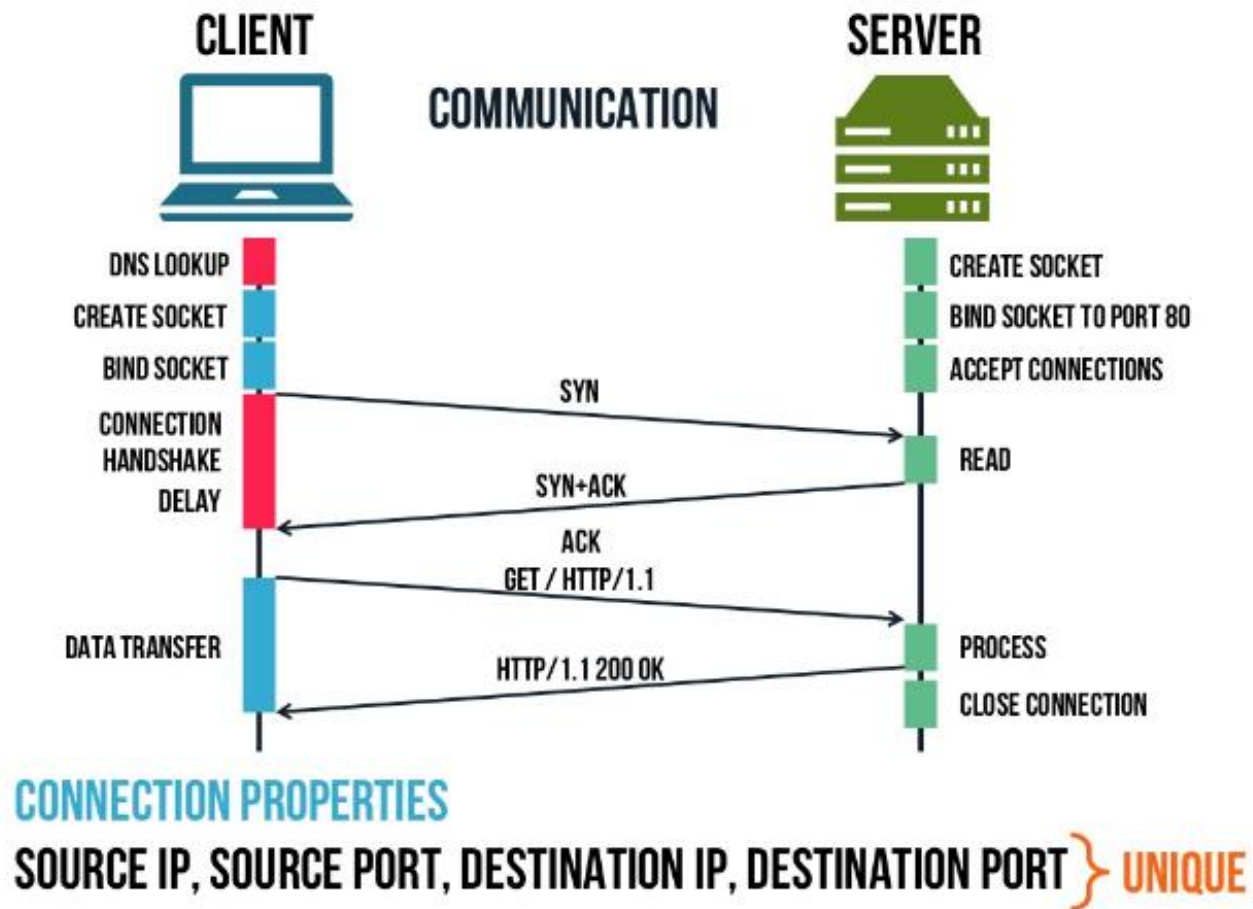
- requested document not found on this server

505 HTTP Version Not Supported



Can you write an
HTTP protocol
sequence diagram?

Can you write an
HTTP protocol
sequence diagram?



HTTP overview

uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains no information about past client requests

aside


protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

■ Non-Persistent HTTP vs. Persistent HTTP

What if there
are more than
one objects?

per·sist·ent

/pər'sist(ə)nt/ 

adjective

1. continuing firmly or obstinately in a course of action in spite of difficulty or opposition.
"one of the government's most persistent critics"
synonyms: tenacious, persevering, determined, resolute, purposeful, dogged, single-minded, tireless, indefatigable, patient, unflagging, untiring, insistent, importunate, relentless, unrelenting; More
2. continuing to exist or endure over a prolonged period.
"persistent rain will affect many areas"
synonyms: constant, continuous, continuing, continual, nonstop, never-ending, steady, uninterrupted, unbroken, interminable, incessant, unceasing, endless, unending, perpetual, unremitting, unrelenting, relentless, unrelieved, sustained More

The background of the slide features a series of concentric circles in light gray, some solid and some dashed, creating a ripple effect. A large, light gray speech bubble is positioned on the left side, containing the title text.

Non-Persistent HTTP vs. Persistent HTTP

Example: an HTML file with 2
pictures in it, how many objects?

HTTP connections

non-persistent HTTP

- at most one object sent over TCP connection
 - connection then closed
- downloading multiple objects required multiple connections

persistent HTTP

- multiple objects can be sent over single TCP connection between client, server

The background of the slide features a series of concentric circles in a light gray color, centered on the left side. A gray speech bubble with a tail pointing towards the bottom left is positioned in the center-left area. The text "Non-Persistent HTTP" is written inside this speech bubble.

Non-Persistent
HTTP


Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =

$2\text{RTT} + \text{file transmission time}$



The background of the slide features a series of concentric circles in light gray, centered on the left side. A large, light gray speech bubble with a pointed tail at the bottom is positioned on the left, containing the text 'Persistent HTTP'.

Persistent HTTP

Non-persistent HTTP: response time

non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

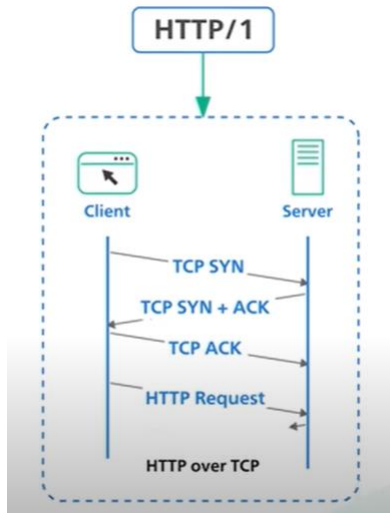
persistent HTTP:

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

HTTP Evolution

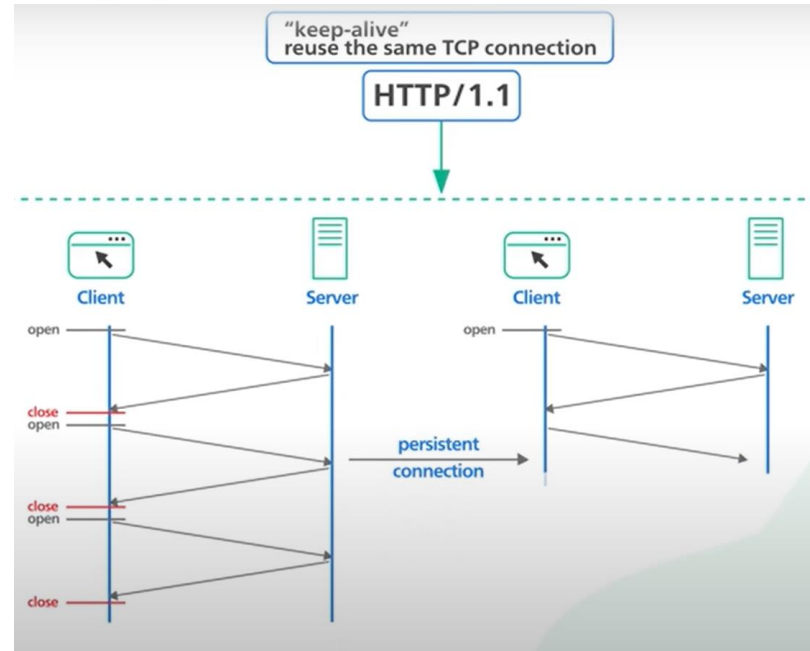
■ HTTP v.1

- HTTP v.1.0




1996

- HTTP v.1.1



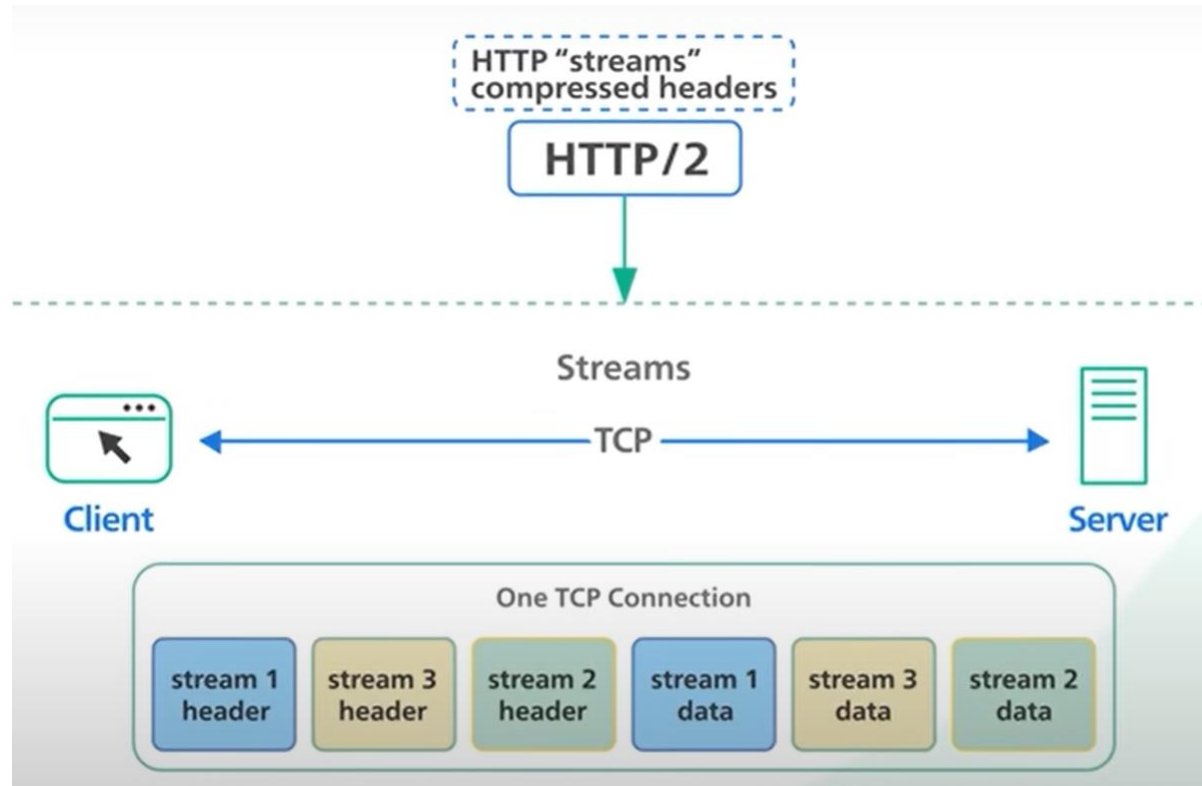
HTTP 1.1 in 1997

- Persistent connection
- Pipelining



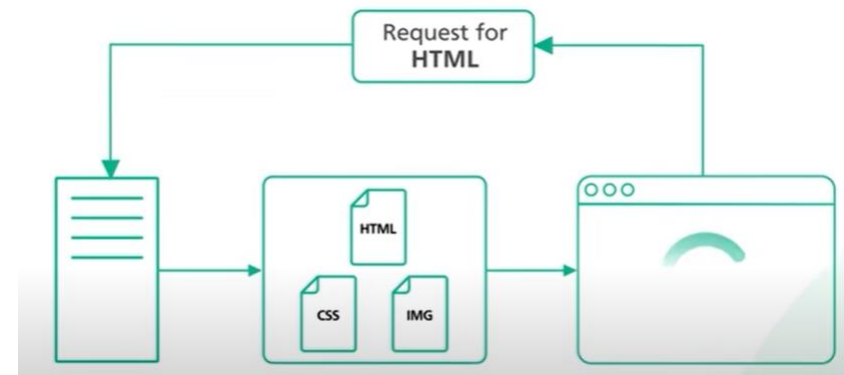
HTTP Evolution

■ HTTP v.2



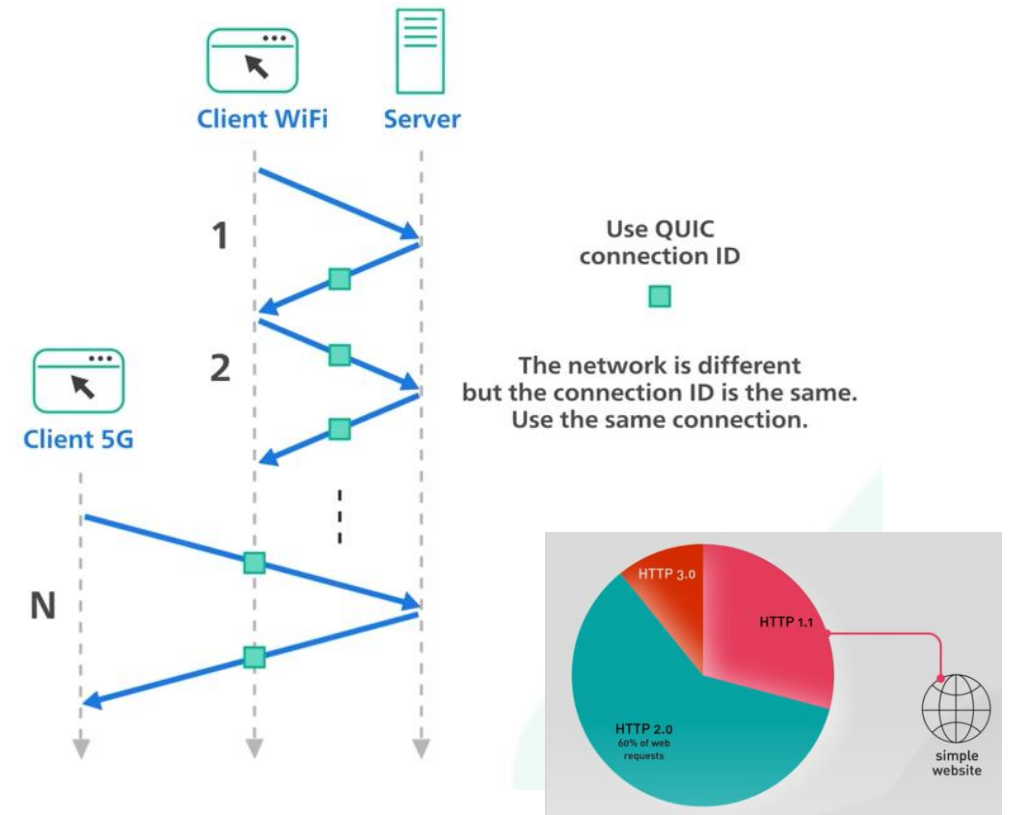
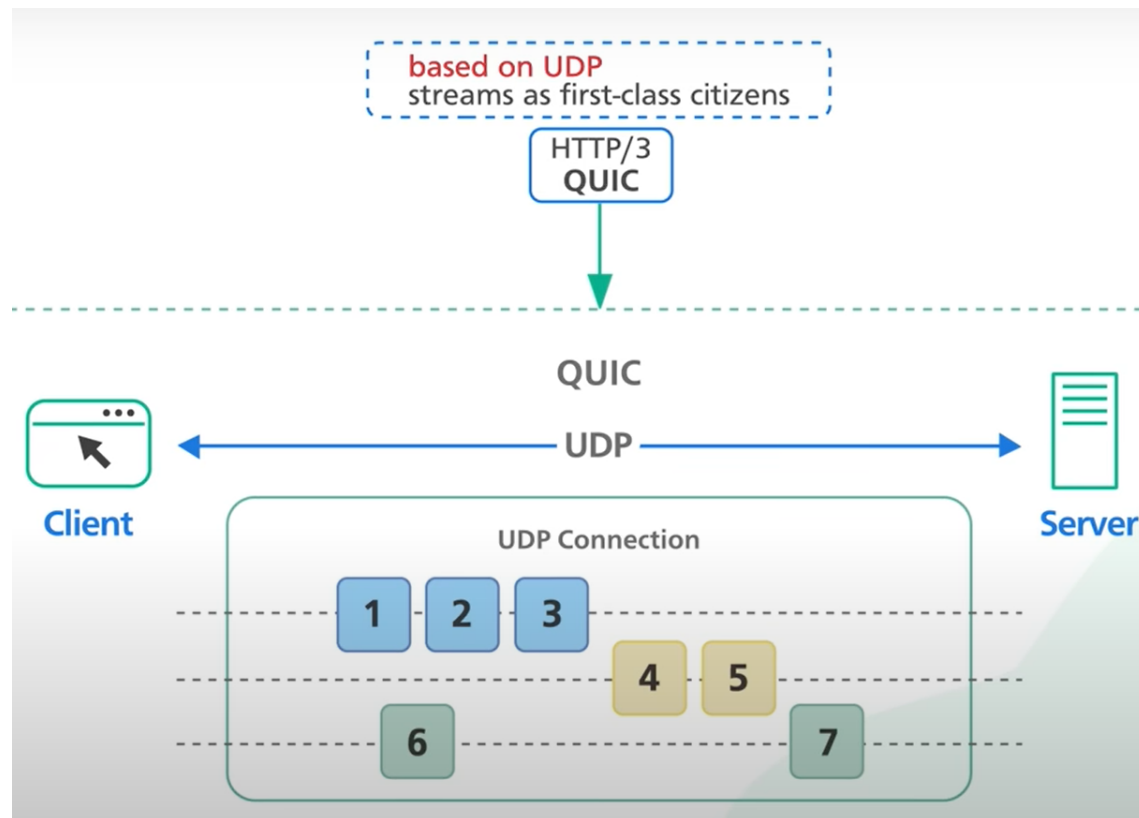
HTTP 2 in 2015

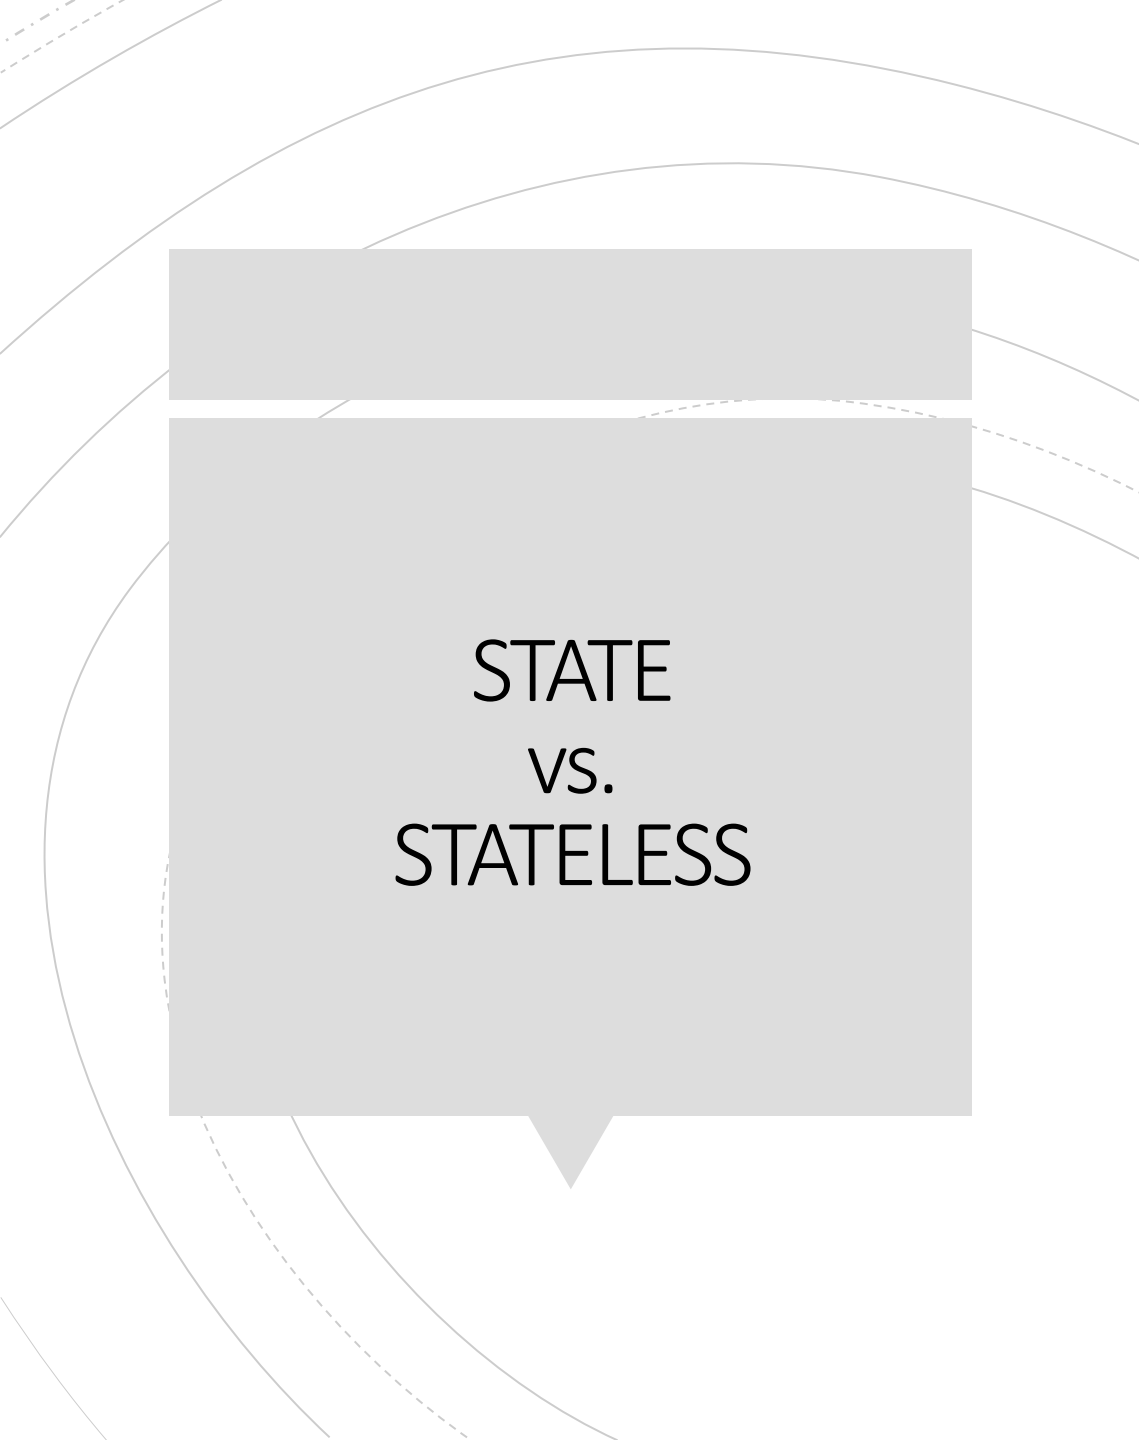
- HTTP stream
- Compressed Header
- Push method introduced : server sends update to client



HTTP Evolution

- **HTTP v.3** HTTP 3 in 2022
 - QUIC protocol





STATE VS. STATELESS

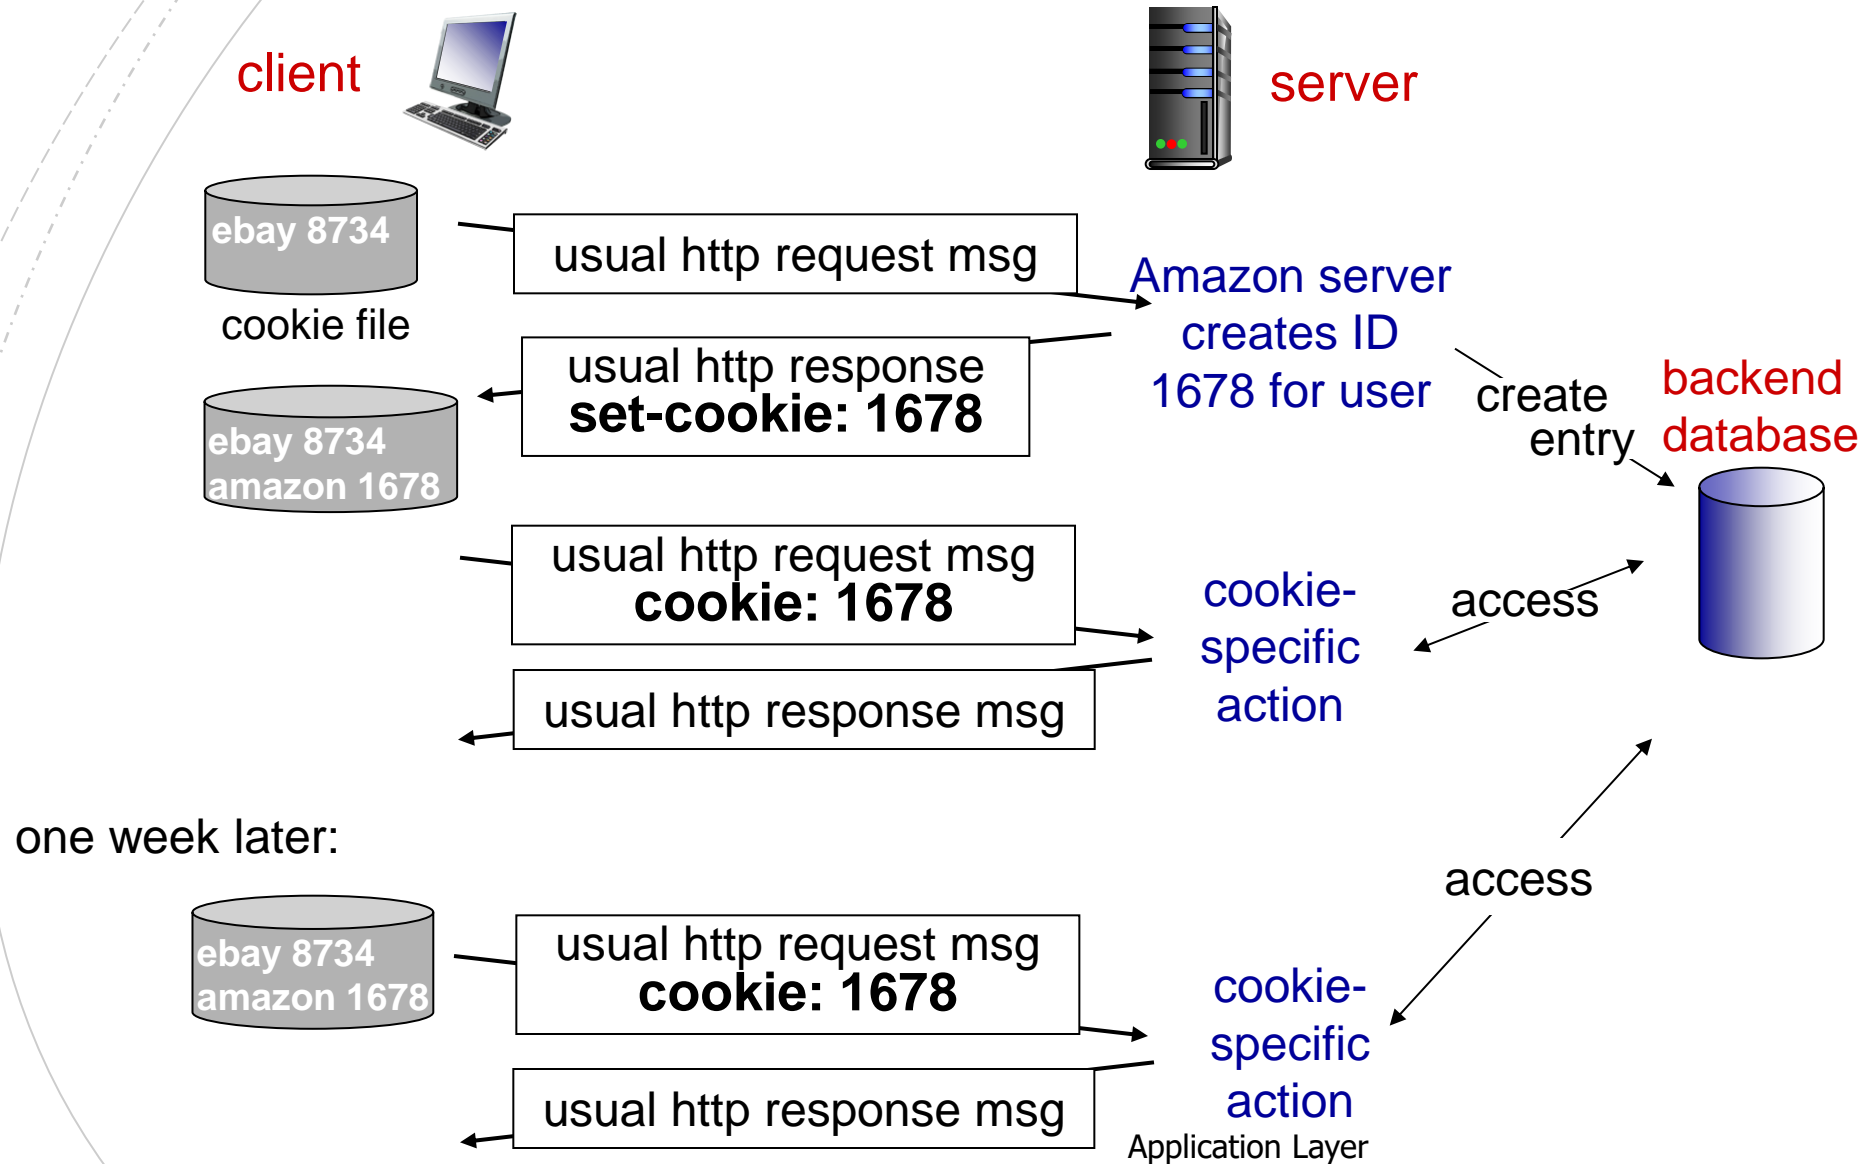
HTTP is stateless,
How to keep track of user information?

HTTP IS STATELESS!!!

We use
cookies to
keep state of
information



Cookies: keeping "state"



User-server state: cookies

many Web sites use cookies

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

example:

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

User-server state: cookies

what cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)


how to keep “state”:

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

aside

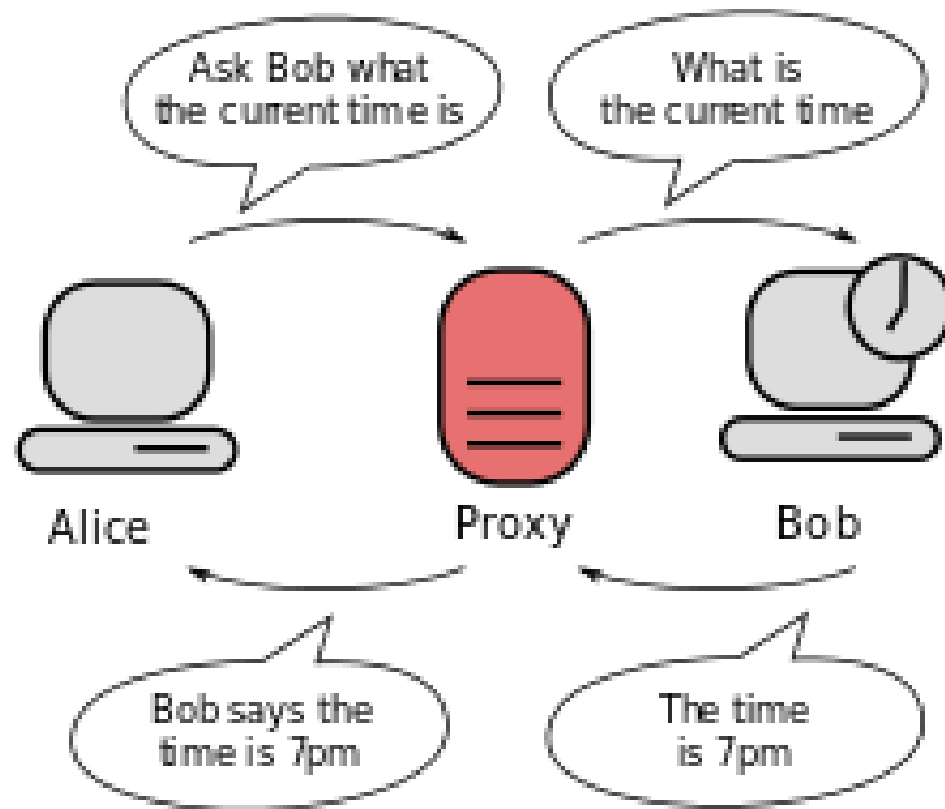
cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

The left side of the slide features a decorative background with several concentric circles in light gray. A large, light gray speech bubble with a tail pointing towards the bottom left is positioned in the center of this section. Inside the speech bubble, the text "Web caches or proxy server" is written in a black, sans-serif font. Above the speech bubble, there is a solid gray rectangular bar.


Web caches
or
proxy server

Can we **FETCH** a web page more efficiently?




Can we FETCH a web page more efficiently?

- **A proxy server may reside on the user's local computer, or at various points between the user's computer and destination servers on the Internet.**
- **A proxy server that passes requests and responses unmodified is usually called a gateway or sometimes a tunneling proxy.**
- - **A forward proxy** is an Internet-facing proxy used to retrieve from a wide range of sources (in most cases anywhere on the Internet).
- - **A reverse proxy** is usually an internal-facing proxy used as a front-end to control and protect access to a server on a private network. A reverse proxy commonly also performs tasks such as load-balancing, authentication, decryption or caching.



An open proxy forwarding requests from and to anywhere on the Internet.




A reverse proxy taking requests from the Internet and forwarding them to servers in an internal network. Those making requests connect to the proxy and may not be aware of the internal network.

Proxy

Web caches (proxy server)

goal: satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



Web caches (proxy server)

- cache acts as both client and server
 - server for original requesting client
 - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

why Web caching?


- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)

Caching example:

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

Can you compute LAN utilization, access link utilization, total delay?



Caching example:


2-44

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = 99% *problem!*
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs




Caching example: fatter access link

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: **154 Mbps**

If we change the access link to 154 Mbps, what will happen?



Caching example: fatter access link


assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: **154 Mbps**

consequences:

- LAN utilization: 15%
- access link utilization = **0.99%**
- total delay = Internet delay + access delay + LAN delay
= **msecs**

HOWEVER INCREASE
ACCESS LINK IS
EXPENSIVE !!!



Caching example: install local cache


assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

*Adding a cache with
a hit-rate of 0.4!*

*How to compute link
utilization, delay?*


Cost: web cache (cheap!)



Caching example: install local cache

Calculating access link utilization, delay with cache:


- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - utilization $= 0.9 / 1.54 = .58$
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
 - less than with 154 Mbps link (and cheaper too!)



Caching example: install local cache

Calculating access link utilization, delay with cache:

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - utilization $= 0.9 / 1.54 = .58$
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
 - less than with 154 Mbps link (and cheaper too!)



Conditional GET

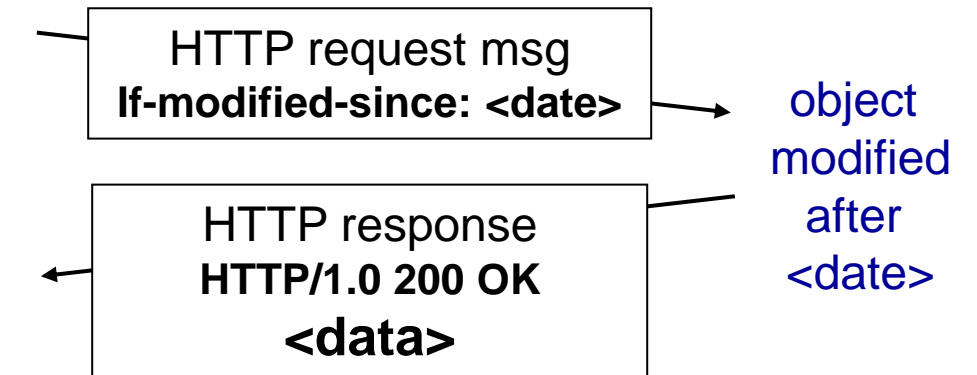
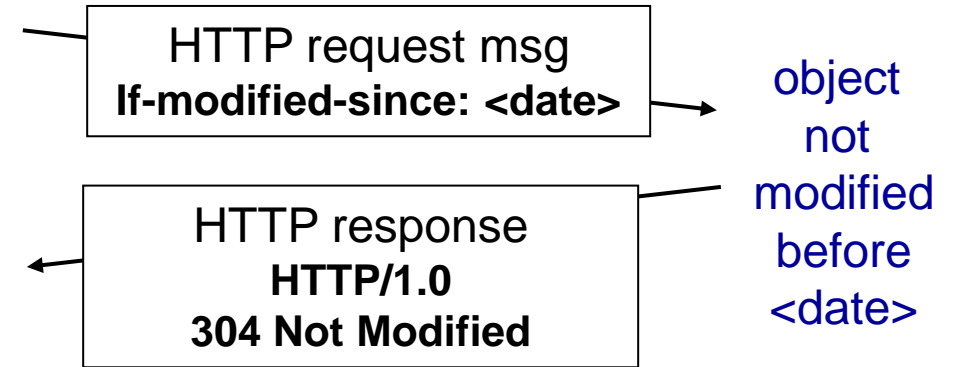
- **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- **cache:** specify date of cached copy in HTTP request

If-modified-since:
<date>
- **server:** response contains no object if cached copy is up-to-date:

HTTP/1.0 304 Not Modified



server



1770	4.073426582	128.119.245.12	10.10.182.50	TCP	74 80 → 55386 [SYN, ACK] Seq=0
1775	4.073525328	10.10.182.50	128.119.245.12	HTTP	613 GET /wireshark-labs/HTTP-wi
2103	4.670905143	128.119.245.12	10.10.182.50	HTTP	305 HTTP/1.1 304 Not Modified

- ▶ Frame 2103: 305 bytes on wire (2440 bits), 305 bytes captured (2440 bits) on interface 0
- ▶ Ethernet II, Src: Industri_29:95:1b (00:00:cd:29:95:1b), Dst: Micro-St_26:88:10 (d8:cb:8a:26:88:10)
- ▶ Internet Protocol Version 4, Src: 128.119.245.12, Dst: 10.10.182.50
- ▶ Transmission Control Protocol, Src Port: 80, Dst Port: 55382, Seq: 1, Ack: 548, Len: 239

▼ Hypertext Transfer Protocol

▶ HTTP/1.1 304 Not Modified\r\n

Date: Mon, 16 Oct 2017 10:21:57 GMT\r\n

Server: Apache/2.4.6 (CentOS) OpenSSL/1.0.2k-fips PHP/5.4.16 mod_perl/2.0.10 Perl/v5.16.3\r\n

Connection: Keep-Alive\r\n

Keep-Alive: timeout=5, max=100\r\n

ETag: "80-55ba3b4031322"\r\n

\r\n

[HTTP response 1/1]

[Time since request: 0.597379815 seconds]

[Request in frame: 1775]